

Mike Wu

(858) 740-9967

wumike@stanford.edu

<https://www.mikehwu.com>

Education	Stanford University 2017 - Present Ph.D. in Computer Science Advisor: Noah Goodman TA: CS398 (Computational Education), CS224S (Spoken Language Processing) Select Courses: CS228 (Probabilistic Graphical Models), CS428 (Computation and Cognition), CS242 (Programming Languages), Math120 (Group Theory), Math115 (Functions of a Real Variable), Math171 (Fundamental Concepts of Analysis), Math175 (Elementary Functional Analysis), Math205a (Measure Theory)
	Yale University , Hopper College 2012 - 2016 B.S. with Distinction in Computer Science Yale college council, science committee Select Courses: CPSC202 (Discrete Math), CPSC223 (Data Structures and Algorithms), CPSC323 (Systems Programming and Computer Arch.), CPSC437 (Databases), CPSC467 (Cryptography and Computer Security)
	University of Oxford , New College 2015 1st mark in three courses in Computer Science Oxford computing society Select Courses: Advanced Data Structures and Algorithms, Machine Learning
Awards and Honors	Pear VC Fellow 2020
	Educational Data Mining Best Paper Award 2020
	Educational Data Mining Best Student Paper Runner-up 2020
	AAAI Outstanding Student Paper Award 2019
	IDEO CoLab Fellow 2019
	AngelHack, Augmented Reality Category 1st place 2018
	AMIA CRI Nominated for Informatics Award 2017
	API World Hackathon, Telesign API 1st place 2017
	Trueface.ai Hackathon, 1st place 2017
	HackMIT Top 8 Hacks, Dropbox API 1st place 2015
	Qualcomm QLiving Scholarship 2014
	Intel ISEF Semifinalist 2012
	Siemens Competition Semifinalist 2012
	Intel ISEF Finalist, 3rd place in Computer Science 2011
	XSEDE Best Student Poster 2011
Fellowships	Stanford Interdisciplinary Graduate Fellowship (Karr Family Fellow) 2020-Current
	NSF Graduate Research Fellowship 2017-2020
Preprints	Mike Wu and Noah Goodman. A Simple Framework for Uncertainty in Contrastive Learning. In <i>ArXiv</i> .
	Mike Wu , Noah Goodman. Multimodal Generative Models for Compositional Rep-

resentation Learning. In *ArXiv*, 2019. Submitted to *Journal of Machine Learning Research* (JMLR).

Ali Malik, **Mike Wu**, Vrinda Vasavada, Jinpeng Song, John Mitchell, Noah Goodman, Chris Piech. Generative Grading: Neural Approximate Parsing for Automated Student Feedback. In *ArXiv*, 2019.

Mike Wu, Sonali Parbhoo, Michael C. Hughes, Ryan Kindle, Leo Celi, Maurizio Zazzi, Volker Roth, Finale Doshi-Velez. Regional Tree Regularization for Interpretability in Black Box Models. In *ArXiv*, 2019. Submitted to *Journal of Artificial Intelligence Research* (JAIR).

**Conference
and Journal
Proceedings**

Alex Tamkin, **Mike Wu**, and Noah Goodman. Viewmaker Networks: Learning Views for Unsupervised Representation Learning. In *Proc. 9th International Conference on Learning Representations* (ICLR), 2021.

Mike Wu, Chengxu Zhuang, Milan Mosse, Daniel Yamins, and Noah Goodman. Conditional Negative Sampling for Contrastive Learning of Visual Representations. In *Proc. 9th International Conference on Learning Representations* (ICLR), 2021.

Mike Wu, Jonathan Nafziger, Anthony Scodary, Andrew Maas. HarperValleyBank: A Domain-Specific Spoken Dialog Corpus. In *ArXiv*, 2020.

Mike Wu, Chengxu Zhuang, Milan Mosse, Daniel Yamins, Noah Goodman. On Mutual Information in Contrastive Learning for Visual Representations. In *ArXiv*, 2020.

Mike Wu, Richard L. Davis, Benjamin W. Domingue, Chris Piech, Noah Goodman. Variational Item Response Theory: Fast, Accurate, and Expressive. In *Educational Data Mining* (EDM), 2020. **[Oral Presentation (15 min).]** **[Best Paper Award.]**

Mike Wu, Sonali Parbhoo, Michael C. Hughes, Volker Roth, Finale Doshi-Velez. Optimizing for Interpretability in Deep Neural Networks with Simulable Decision Trees. In *Proc. 34rd AAAI Conference on Artificial Intelligence* (AAAI), 2020. **[Oral Presentation (20 min).]**

Mike Wu, Kristy Choi, Noah Goodman, Stefano Ermon. Meta-Amortized Variational Inference and Learning. In *Proc. 34rd AAAI Conference on Artificial Intelligence* (AAAI), 2020.

Judith Fan, Robert X.D. Hawkins, **Mike Wu**, Noah Goodman. Pragmatic inference and Visual Abstraction Enable Contextual Flexibility during Visual Communication. In *Computational Brain and Behavior* (COBB), 2019.

Mike Wu, Noah Goodman, Stefano Ermon. Differentiable Antithetic Sampling for Variance Reduction in Stochastic Variational Inference. In *Proc. 22nd International Conference on Artificial Intelligence and Statistics* (AISTATS), 2019.

Mike Wu, Milan Mosse, Noah Goodman, Chris Piech. Zero Shot Learning for Code Education: Rubric Sampling with Deep Learning Inference. In *Proc. 33rd AAAI Conference on Artificial Intelligence* (AAAI), 2019. **[Oral Presentation (12 min).]** **[Outstanding Student Paper Award.]**

Mike Wu, Noah Goodman. Multimodal Generative Models for Scalable Weakly-

Supervised Learning. *Proc. 32nd Annual Conference on Neural Information Processing Systems* (NeurIPS), 2018.

Mike Wu, Michael C. Hughes, Sonali Parbhoo, Maurizio Zazzi, Volker Roth, Finale Doshi-Velez. Beyond Sparsity: Tree Regularization of Deep Models for Interpretability. In *Proc. 32nd AAAI Conference on Artificial Intelligence* (AAAI), 2018. [**Spotlight Presentation (2 min).**]

Marzyeh Ghassemi, **Mike Wu**, Michael C. Hughes, Finale Doshi-Velez. Predicting Intervention Onset in the ICU with Switching Statespace Models. In *Proc. AMIA Summit on Clinical Research Informatics* (CRI), 2017. [**Nominated for Informatics Award.**]

Mike Wu, Marzyeh Ghassemi, Mengling Feng, Leo Anthony Celi, Peter Szolovitz, Finale Doshi-Velez. Understanding Vassopressor Intervention and Weaning: Risk Prediction in a Public Heterogeneous Clinical Time Series Database. In *Journal of the American Medical Informations Association, Volume 24, Issue 3, No. 1* (JAMIA), 2016.

Mike Wu, Madhu Krishnan. Edge-based Crowd Detection from Single Image Datasets. In *International Journal of Computer Science Issues, Volume 12, Issue 1, No. 1* (IJCSI), 2013.

Madhu Krishnan, **Mike Wu**, Young Kang, Sarah H. Lee. Autonomous Mapping and Navigation through Utilization of Edge-based Optical Flow and Time-to-Collision. In *ARPN Journal of Engineering and Applied Sciences, Volume 7, No. 12*, 2012.

Workshops

Mike Wu, Kristy Choi, Noah Goodman, Stefano Ermon. Meta-Amortized Variational Inference and Learning. *NeurIPS 2019 Workshop on Bayesian Deep Learning*. [**Spotlight Presentation (1 min).**]

Mike Wu, Sonali Parbhoo, Finale Doshi-Velez. Beyond Sparsity: Tree Regularization of Deep Models for Interpretability. *NeurIPS 2017 Workshop on Transparent and interpretable Machine Learning in Safety Critical Environments*. [**Contributed Talk (10 min).**]

Patents

Frank Wood, **Mike Wu**, Yura Perov, Hongseok Yang. Computing engine, software, system and method. US Patent App. 15/465,131, 2017.

Teaching Experience

Teaching Assistant, Dept. of Computer Science, Stanford University Spring 2020 CS224S: Spoken Language Processing (Andrew Maas)

Teaching Assistant, Dept. of Computer Science, Stanford University Fall 2019 CS398: Computational Education (Chris Piech)

Teaching Assistant, Dept. of Computer Science, Yale University Spring 2016 CPSC437: Operating System Concepts (Avi Silberschatz)

Teaching Assistant, School of Management, Yale University Fall 2015 MGT656: Management of Software Development (Kyle Jensen)

Invited Talks

Multidisciplinary University Research Initiative (MURI) Visual Common-Sense Annual Meeting, 2020.

Berkeley Center for Human-Compatible AI Technical Seminar, 2020.

Stanford Computer Forum Annual Meeting, 2019.

Human-Centered Artificial Intelligence Institute Symposium, 2019.

Yale Technology Conference, Yale University, 2016.

Probabilistic Programming Workshop, University of Southampton, 2016.

Conference Abstracts

Judith Fan, Robert X.D. Hawkins, **Mike Wu**, Noah Goodman. Modeling contextual flexibility in visual communication. *Vision Sciences Society Annual Meeting (VSS)*, 2018.

William Smith, **Mike Wu**, Yura Perov, Frank Wood, Hongseok Yang. Spreadsheet probabilistic programming. *Inaugural International Conference on Probabilistic Programming*. (PROBPROG), 2018.

Jessi Cisewski-Kehe, **Mike Wu**, Brittany Fasy, Wojciech Hellwing, Mark Lovell, Alessandro Rinaldo, Larry Wasserman. Investigating the Cosmic Web with Topological Data Analysis. *American Astronomical Society Meeting*. (AAS) 2018.

Industry Experience

Foundation Capital	2020-Current
Internship with Li Sun.	

Facebook Applied Machine Learning (AML)	2016-2017
Visiting engineer building tools in computer vision.	

Lattice Data	2016
Software engineer building weakly supervised classifiers.	

Invrea (Inverse Reasoning) http://invrea.com	2015-2017
Co-founder building a probabilistic programming language in Excel.	

YHack https://www.yhack.org	2013-2016
Co-founder building an international collegiate hackathon.	

Ionis Pharmaceuticals	2013
Data science Intern building classifiers for drug design.	